

# Crossing Language Barriers with julia, SciPy, and IP[y]thon

Steven G. Johnson  
MIT Applied Mathematics

# Where I'm coming from...

[ google "Steven Johnson MIT" ]

Computational software you may know...

... mainly C/C++ libraries & software ...

... often with Python interfaces ...


(& Matlab & Scheme & ...)

Nanophotonics




[www.fftw.org](http://www.fftw.org)


[jdj.mit.edu/nlopt](http://jdj.mit.edu/nlopt)



$\text{erf}(z)$  (and  $\text{erfc}$ ,  $\text{erfi}$ , ...)  
in SciPy 0.12+



& other EM simulators...




[jdj.mit.edu/book](http://jdj.mit.edu/book)

**Confession:** I've used Python's internal C API more than I've coded in Python...

# A new programming language?

Jeff Bezanson



Viral Shah



Stefan Karpinski

Alan Edelman



[ 17+ developers with 100+ commits ]



[julialang.org](http://julialang.org)

[begun 2009, “0.1” in 2013, ~20k commits]

First reaction: **You're doomed.** [ usual fate of all new languages ]

... subsequently:

... ~~probably doomed~~  
... still might be doomed

but, in the meantime,

I'm **having fun** with it...

... and **it solves a real problem**  
with technical computing  
in high-level languages.

# The “Two-Language” Problem

Want a **high-level language** that you can work with **interactively**  
= easy development, prototyping, exploration  
⇒ **dynamically typed language**

*Plenty to choose from: Python, Matlab / Octave, R, Scilab, ...*  
(& some of us even like Scheme / Guile)

Historically, **can't write performance-critical code** (“inner loops”)  
in these languages... **have to switch to C/Fortran/...** (static).  
[ e.g. SciPy git master is ~70% C/C++/Fortran]

Workable, but **Python → Python+C = a huge jump in complexity.**

# Just vectorize your code?

= rely on mature **external libraries**,  
operating on **large blocks of data**,  
for performance-critical code

**Good advice!** But...

- **Someone** has to write those libraries.
- Eventually that person may be **you**.
  - **some problems** are impossible or just very awkward to vectorize.

# Dynamic languages don't have to be slow.

Lots of progress in JIT compilers, driven by web applications.  
& excellent free/open-source JIT via LLVM.

Javascript in modern browsers achieves C-like speed.

Many other efforts to speed up dynamic languages, e.g. PyPy,  
Numba / Cython (really 2<sup>nd</sup> lower-level language embedded in Python).

*What if a dynamic language were designed for JIT  
from the beginning, with the goal of being as high-  
level as possible while staying within  $2\times$  C speed?*



(and it's easier to call SciPy  
from Julia than from PyPy)



# *Today*

A **brief introduction to Julia**,  
its key features,  
and how it gets performance.

How **Julia leverages Python and IPython**  
to lessen the “infrastructure problem” of new languages

*time permitting:*

How tools can **flow in the other direction** too...



[ [julialang.org](http://julialang.org) ]

- Dynamically typed
- Multiple dispatch: a generalization of OO
- Metaprogramming (code that writes code)
- Direct calling of C and Python functions
- Coroutines, asynchronous I/O
- Designed for Unicode
- Distributed-memory parallelism
- User-defined types == builtin types ...  
extensible promotion and conversion  
rules, etc.
- Large built-in library: regex, linear algebra,  
special functions, integration, etcetera...
- git-based package manager

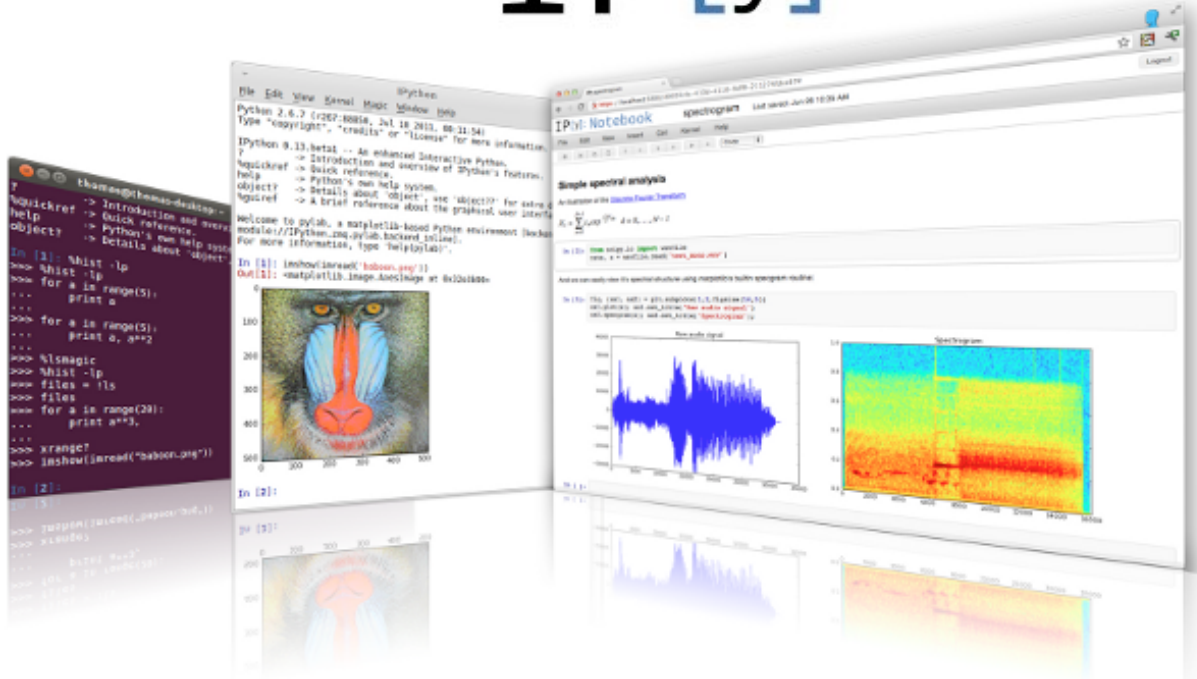
Most of Julia (70%+)  
is written in Julia.

goto live Julia REPL demo...

# This command line is so 1990...

[e.g. GNU Readline, 1989]

vs. **IP[y]**



Come back in 5 years  
once you implement  
something more  
modern...

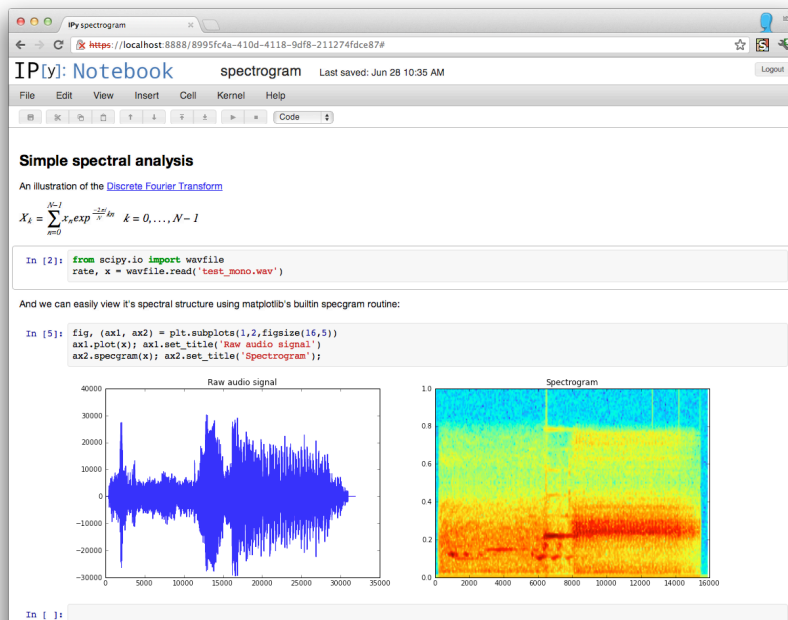
[ [ipython.org](http://ipython.org) ]

# (roughly) How IPython Notebooks Work

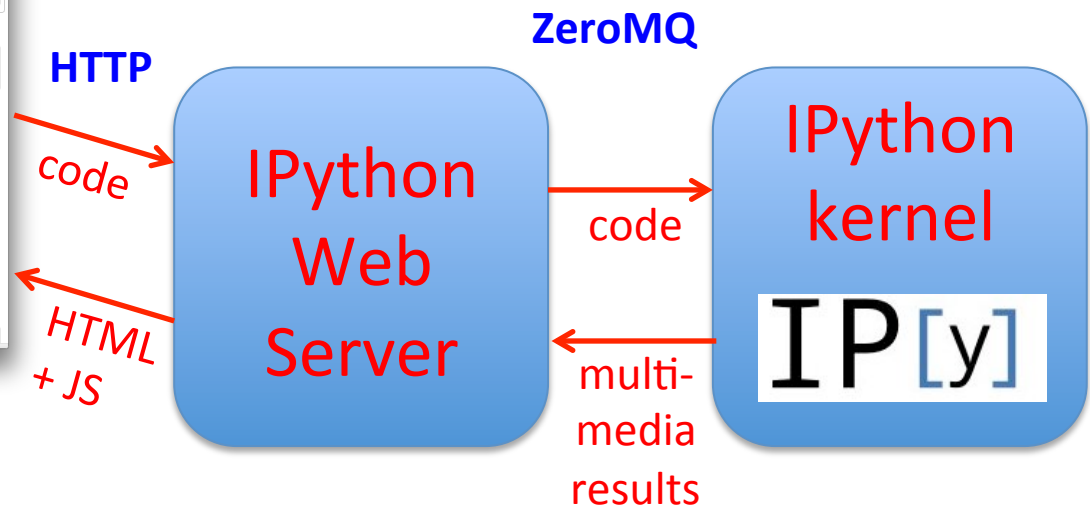
programmer



enter Python code



Web browser:  
Notebook Display/Interaction

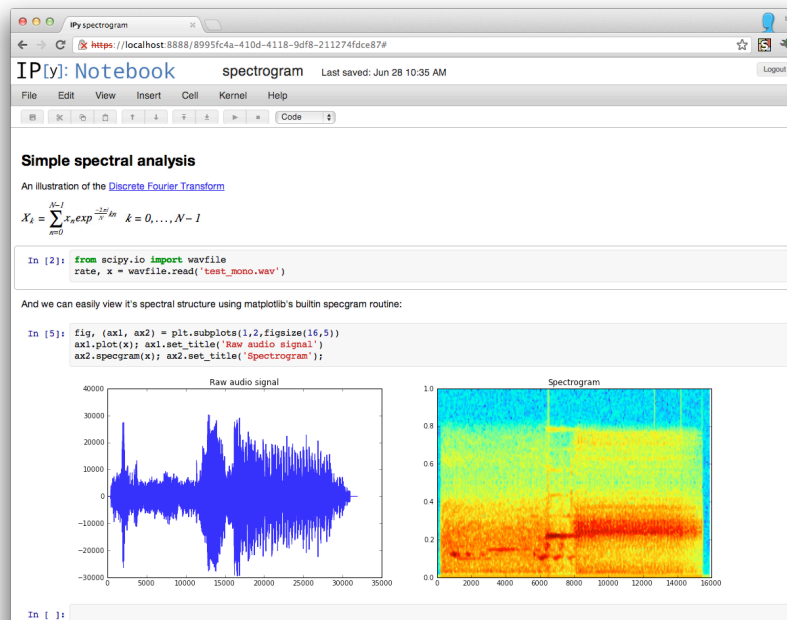


# How Julia Notebooks Work

programmer



enter Julia code



Web browser:

Notebook Display/Interaction

HTTP

code

HTML  
+ JS

ZeroMQ

~~Python~~  
Jupyter  
Web  
Server

code

multi-  
media  
results

Julia  
kernel



goto live Julia notebook demo...


Why **is** Julia fast?



~~Why~~ is Julia fast?

# Julia performance on synthetic benchmarks

[ loops, recursion, etc., implemented in most straightforward style ]



What about **real problems**,  
compared to **highly optimized** code?

# Special Functions in Julia

Special functions  $s(x)$ : classic case that cannot be vectorized well

... switch between various polynomials depending on  $x$

Many of Julia's special functions come from the usual C/Fortran libraries, but **some** are written in **pure Julia** code.

Pure Julia **erfinv**( $x$ ) [ =  $\text{erf}^{-1}(x)$  ]

**3–4× faster than Matlab's** and **2–3× faster than SciPy's** (Fortran Cephes).

Pure Julia **polygamma**( $m, z$ ) [ =  $(m+1)^{\text{th}}$  derivative of the  $\ln \Gamma$  function ]

**~ 2× faster than SciPy's** (C/Fortran) for real  $z$


... and unlike SciPy's, *same code* supports complex argument  $z$

Julia code can actually be **faster** than typical “optimized” C/Fortran code, by using **techniques** [metaprogramming/**code generation**] that are **hard in a low-level language**.

# Pure-Julia FFT performance

double-precision complex, 1d transforms

powers of two



already comparable to FFTPACK

[ probably some tweaks to  
inlining will make it better ]

FFTW 1.0-like code generation  
+ recursion in Julia  
~ 1/3 lines of code compared to  
FFTPACK, more functionality

Why **is** Julia fast?

~~Why **is** Julia fast?~~

Why **can** Julia be fast?

(You can write slow code in any language, of course.)

... and couldn't Python do the same thing?

# Type Inference

To generate **fast code for a function  $f(x,y)$** , the compiler needs to be able to **infer the types of variables in  $f$** , map them to hardware types (registers) where possible, and **call specialized code paths** for those types (e.g. you want to inline  $+$ , but this depends on types).

At compile-time, the compiler generally **only knows types of  $x,y$ , not values**, and it needs to be able to cascade this information to infer types throughout  $f$  and in any functions called by  $f$ .

**Julia and its standard library are designed so type inference is possible** for code following straightforward rules.

... sometimes this requires **subtle choices** that would be **painful to retrofit** onto an existing language.



# Type Stability

Key to predictable, understandable type inference:

- the **type of function's return value** should **only depend** on the **types of its arguments**

A counter-example in Matlab and GNU Octave:

`sqrt(1) == 1.0` — real floating-point  
`sqrt(-1) == 0.0+1.0i` — complex floating-point

Hence, any **non-vector code that calls `sqrt(x)`** in Matlab **cannot be compiled** to fast code **even if `x` is known to be real scalar** — anything “touched” by the `sqrt(x)` is “poisoned” with an unknown type — unless the compiler knows  $x \geq 0$ .

Better to throw an exception for `sqrt(-1)`, requiring `sqrt(-1+0i)`.

# Type Stability

Key to predictable, understandable type-inference:

- the type of function's return value should only depend on the types of its arguments

Common counter-examples in Python

Typical idiom:

foo(x) returns **y, or None** if [exceptional condition]

[e.g. `numpy.ma.notmasked_edges`, `scipy.constants.find`, ...]

Better: **Throw an exception.**

# Type Stability

Key to predictable, understandable type-inference:

- the type of function's return value should only depend on the types of its arguments

A counter-example in Python

integer arithmetic

Integer arithmetic in Python automatically uses bignums to prevent overflow. Unless the compiler can detect that overflow is impossible [which **may be detectable sometimes!**], integers can't be compiled to integer registers & hw arithmetic.

**Julia tradeoff:** default **integers are 64-bit**, overflow possible  
... use explicit BigInt type if you are doing number theory.

goto live Julia notebook demo...

**Julia:** fun, fast, and  
you don't lose your Python stuff.

New languages are always a risk...

...but maybe not doomed?

# Acknowledgements



[julialang.org](http://julialang.org)

Julia core team:

Jeff Bezanson (MIT)

Stefan Karpinski (MIT)

Viral Shah

*...(17+ developers with 100+ commits)...*

Prof. Alan Edelman (MIT)



IP[y]

[ipython.org](http://ipython.org)



Fernando  
Perez



Matthias  
Bussonier



Min RK

& Shashi  
Gowda  
(GSoC)